

Haiyang Yu

<https://www.linkedin.com/in/haiyang-yu-246a213b/>

haiyangy@andrew.cmu.edu

+1-408-839-3245

Cupertino, CA

EDUCATION

- **Carnegie Mellon University Silicon Valley** Mountain View
Master of Science in Software Engineering Jan. 2018 – May. 2019
- **Nanyang Technological University** Singapore
Bachelor of Engineering, Electrical and Electronic Engineering; GPA: 4.61/5.0 Aug. 2010 – May. 2014
- **Courses:** Introduction to Machine Learning for Engineers, Practical Data Science, Foundations of Computer Systems

TECHNICAL SKILLS

- **Languages:** Python, Java, C, Shell, JavaScript, SQL
- **Technologies:** NumPy, Pandas, scikit-learn, Docker, Linux

PROJECTS

- **Movie Box Office Prediction (Python):** Crawled raw html of over 8,000 movies from Box Office Mojo and parsed them into meaningful data using **lxml**. Processed 30 million pieces of movie information using **MapReduce**. Extracted features using **Pandas**, **NumPy** and **One Hot Encoding**. Compared prediction accuracy between models trained using **Linear Regression** and **Decision Tree Regression**.

WORK EXPERIENCE

- **Moqi.ai (Java, Shell)** Beijing, China
Software Engineer Intern May 2018 - Aug 2018
Designed and implemented fingerprint system fault tolerance that improved system uptime from 93% to 99%
 - **Distributed Storage:** Set up **Ceph Filesystem** with **Erasur Coding** on SSD as persistent distributed storage for segments cached in individual matching servers' memory.
 - **Fault Tolerance:** Designed and implemented auto discovery of online/offline **dockerized** matching servers through **Consul** and **Registrar**, which triggers segment redistribution among matching servers and segment recovery from Ceph Filesystem.
 - **Docker Compose:** Simplified fingerprint system testing through creating **docker compose** file to run the entire fingerprint system on a single server. The fingerprint system consists of heterogeneous **gRPC** services written in Java, Python, C++ as well as Redis, Cassandra and SeaweedFS.
- **Barclays Capital Services (Java)** Singapore
Software Engineer Jun 2014 - Dec 2017
 - **Trade Reporting Processor:** Developed Trade Reporting Processor using **Google Guava** to cache stock indices which improved performance by 2x.
 - **Market Object and Static Services:** Developed a greenfield market data calculation application which used **Redis** for caching, **Elasticsearch** for logging and **Akka** for concurrent processing.
 - **Sparta Automatic Deployment:** Fully automated deployment workflow to 24 servers across 5 countries that consisted of migrating Sparta from **Perforce** to **Git** version control system, and implementing continuous deployment through **TeamCity** and **Nolio**.